



上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY

CS075



# Lecture 14

## Final Projects

**Jianwen Wei**

Center for HPC, SJTU

Jun 2019

# The project report for CS075

- Picking one or more projects to finish a report for your final scores of CS075.
- Length: >2 pages. Written in English.
- Due : **23:59 Jul 5<sup>th</sup>, 2019**
- Work in the project should be done independently and can be reproduced.
- Sample codes can be found <https://hpc.sjtu.edu.cn/info/1016/1697.htm> . Code on github can be used in the project report.
- To get a high score, you can try
  - Finishing more than one project.
  - Finishing the "extra credits" part in projects
  - Fixing bugs in the project instructions.
  - Presenting what you learn in this course
- Send your report and feedbacks to TA: **huhang123@sjtu.edu.cn** . Read the docs <https://pi.sjtu.edu.cn/doc/> carefully before sending your issues to TA.

# A project report should contain

- A brief introduction of SJTU π supercomputer.
- Background of your applications and motivation of your work.
- Experiment setups.
- Performance analysis.
- Summary

# Project 1: OpenMP



- Running an OpenMP program to calculate the value of  $\pi$  and comparing performance against the number of cores.
- Reference instructions:
  - <https://computing.llnl.gov/tutorials/openMP/>
  - <http://hpc.sjtu.edu.cn/info/1016/1477.htm>
- Extra credits
  - Comparing performance between cases with and without false sharing between threads.
  - Comparing performance between different compilers
  - Comparing performance between different OpenMP scheduler strategies
  - Running and analyzing performance of OpenMP programs other than  $\pi$  calculation

# Project 2: MPI



- Running an MPI program to calculate the value of  $\pi$  and comparing performance against the number of cores.
- Reference instructions:
  - <https://computing.llnl.gov/tutorials/mpi/>
  - <http://hpc.sjtu.edu.cn/info/1016/1477.htm>
- Extra credits
  - Comparing performance between different compilers and MPI suites
  - Comparing performance between cases where MPI processes within a same node and among different nodes.
  - Running and analyzing performance of MPI programs other than  $\pi$  calculation

# Project 3: RDMA



- Running a RDMA program two nodes and evaluating RDMA performance in messages per seconds v.s. cpu utiliation.
- Reference instructions:
  - <https://community.mellanox.com/docs/DOC-2493>
  - [http://www.mellanox.com/related-docs/prod\\_software/RDMA\\_Aware\\_Programming\\_user\\_manual.pdf](http://www.mellanox.com/related-docs/prod_software/RDMA_Aware_Programming_user_manual.pdf)
  - [https://github.com/efficient/rdma\\_bench](https://github.com/efficient/rdma_bench)
- Extra credits
  - Communicating between more than 2 nodes.
  - Comparing performance between TCP/IP and RDMA when doing similar operations.
  - Trying more than one RDMA protocol and comparing their performance and programmability.
  - Running a real application (for example, scp) on RDMA and comparing performance between IP and RDMA.
  - Running RDMA Benchmarks

# Project 4: Apache Spark



- Starting a two-node Spark cluster then running SparkPi program to calculate  $\pi$ .
- Reference instructions:
  - <https://pi.sjtu.edu.cn/doc/spark/>
  - <https://github.com/LLNL/magpie>
- Extra credits
  - Running other Spark programs.
  - Running other BigData Analysis stacks on Spark
  - Launching Spark programs via Magpie scripts
  - Using RDMA (for example IPolB) for inter-node communication

# Project 5: CUDA



- Running the CUDA code to calculate the matrix-matrix multiplication and compare the performance with the sequential code to get the speedup.
- Reference instructions:
  - <http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>
  - <https://pi.sjtu.edu.cn/doc/samples/>
- Extra credits
  - Optimizing the CUDA code using shared memory.
  - Optimizing the matrix-matrix multiplication algorithm with blocking.
  - Could the input matrix be with any size not only the square matrix?

# Project 6: OpenACC



- Add OpenACC directives to the matrix-matrix multiplication C code and run the OpenACC with PGI compiler on GPU card.
- Reference instructions:
  - [http://www.openacc.org/sites/default/files/inline-files/OpenACC\\_Programming\\_Guide\\_0.pdf](http://www.openacc.org/sites/default/files/inline-files/OpenACC_Programming_Guide_0.pdf)
  - <https://pi.sjtu.edu.cn/doc/modules/>
- Extra credits
  - Comparing performance with different compilers options.
  - Try different gang, vector size to achieve the best performance on GPU.
  - Running the OpenACC code on x86 multicore CPU.

# Project 7: PyTorch CNN



- Run a CNN program on PyTorch.
- Reference Instructions
  - <https://github.com/MorvanZhou/PyTorch-Tutorial>
- Extra credits
  - Comparing performance between CPU and GPU.
  - Comparing performance between different settings of hyper parameters.

# Project 8: PyTorch RNN



- Run a RNN program on PyTorch.
- Extra credits
  - Comparing performance between CPU and GPU.
  - Comparing performance between different settings of hyper parameters.